

# Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

**Objective:**

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## **[1]. Reading Data**

## [1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [2]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
```

```

from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

```

```

C:\Users\nairh\Anaconda3\lib\site-packages\smart_open\ssh.py:34: UserWarning: paramiko missing, opening SSH/SCP/SFTP paths will be disabled. `pip install paramiko` to suppress
  warnings.warn('paramiko missing, opening SSH/SCP/SFTP paths will be disabled. `pip install paramiko` to suppress')
C:\Users\nairh\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
  warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

```

In [3]:

```

# using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3

```

```

# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (100000, 10)

Out[3]:

Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0					
1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1

1

2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0
---	------------	----------------	--------	---

2

3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1
---	------------	---------------	--	---



In [4]:

```
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COU
NT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

In [5]:

```
print(display.shape)
display.head()
```

(80668, 7)

Out[5]:

	UserId	ProductId	ProfileName	Time	Score	Text
0	#oc- R115TNMSPFT9I7	B005ZBZLT4	Breyton	1331510400	2	Overall its just OK when considering the price...
1	#oc- R11D9D7SHXIJB9	B005HG9ESG	Louis E. Emory	1342396800	5	My wife has recurring extreme

"hoppy"

muscle  
spasms, u...

2	#oc- R11DNU2NBKQ23Z	B005ZBZLT4	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...
3	#oc- R11O5J5ZVQE25C	B005HG9ESG	Penguin Chick	1346889600	5	This will be the bottle that you grab from the..
4	#oc- R12KPBODL2B5ZD	B007OSBEV0	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...




In [6]:

```
display[display['UserId']== 'AZY10LLTJ71NX']
```

Out[6]:

	UserId	ProductId	ProfileName	Time	Score	Text
80638	AZY10LLTJ71NX	B001ATMQK2	undertheshrine "undertheshrine"	1296691200	5	I bought this 6 pack because for the price tha...



In [6]:

```
display['COUNT(*)'].sum()
```

Out[6]:

393063



## **[2] Exploratory Data Analysis**

## [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [8]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[8]:

	Id	ProductId	UserId	ProfileName
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan

4

155049 B000PAQ75C AR5J8UI46CURR

Geetha  
Krishnan

2



As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [9]:

```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

In [10]:

```
#Deduplication of entries
```

```
final=sorted_data.drop_duplicates(subset={"UserId", "ProfileName", "Time", "Text"}, keep='first', inplace=False)  
final.shape
```

Out[10]:

```
(87775, 10)
```

In [11]:

```
#Checking to see how much % of data still remains
```

```
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[11]:

```
87.775
```

**Observation:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [12]:

```
display= pd.read_sql_query("""  
SELECT *  
FROM Reviews  
WHERE Score != 3 AND Id=44737 OR Id=64422  
ORDER BY ProductID  
""", con)  
  
display.head()
```

Out[12]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator
0				J. E.	

64422	B000MIDROQ	A161DK06JJMCYF	Stephens "Jeanne"	3
-------	------------	----------------	----------------------	---

1

44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3
-------	------------	----------------	-----	---



In [13]:

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [14]:

```
#Before starting the next phase of preprocessing lets see the  
number of entries left  
print(final.shape)  
  
#How many positive and negative reviews are present in our da  
taset?  
final['Score'].value_counts()
```

(87773, 10)

Out[14]:

```
1    73592  
0    14181  
Name: Score, dtype: int64
```

In [46]:

```
timesorted_data=final.sort_values('Time') #, axis=0, ascending=True,  
inplace=False, kind='quicksort', na_position='last')
```

In [53]:

```
print(timesorted_data.shape)
```

```
(87773, 10)
```

## [3] Preprocessing

In [ ]:

## [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [15]:

```
# printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)
```



```
sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

My dogs loves this chicken but its a product f  
rom China, so we wont be buying it anymore. I  
ts very hard to find any chicken products made  
in the USA but they are out there, but this o  
ne isnt. Its too bad too because its a good p  
roduct but I wont take any chances till they k  
now what is going on with the china imports.

=====  
=====

The Candy Blocks were a nice visual for the Le  
go Birthday party but the candy has little tas  
te to it. Very little of the 2 lbs that I bou  
ght were eaten and I threw the rest away. I w  
ould not buy the candy again.

=====  
=====

was way to hot for my blood, took a bite and d  
id a jig lol

=====  
=====

My dog LOVES these treats. They tend to have a  
very strong fish oil smell. So if you are afr  
aid of the fishy smell, don't get it. But I th  
ink my dog likes it because of the smell. Thes  
e treats are really small in size. They are gr  
eat for training. You can give your dog severa  
l of these without worrying about him over eat  
ing. Amazon's price was much more reasonable t  
han any other retailer. You can buy a 1 pound  
bag on Amazon for almost the same price as a 6  
ounce bag at other retailers. It's definitely  
worth it to buy a big bag if your dog eats th

em a lot.

```
=====
=====
```

In [16]:

```
# remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_1500 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.

In [17]:

```
# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)
```

```

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.

=====

====

The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste to it. Very little of the 2 lbs that I bought were eaten and I threw the rest away. I would not buy the candy again.

=====

====

was way to hot for my blood, took a bite and did a jig lol

=====

====

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid of the fishy smell, don't get it. But I think my dog likes it because of the smell. These treats are really small in size. They are great for training. You can give your dog several of these without worrying about him over eat

ing. Amazon's price was much more reasonable than any other retailer. You can buy a 1 pound bag on Amazon for almost the same price as a 6 ounce bag at other retailers. It's definitely worth it to buy a big bag if your dog eats them a lot.

In [18]:

```
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

In [19]:

```
sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

was way to hot for my blood, took a bite and did a jig lol

```
=====
=====
```

In [20]:

```
#remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.

In [21]:

```
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

was way to hot for my blood took a bite and did a jig lol

In [22]:

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have removed in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've",\
```

```

        "you'll", "you'd", 'your', 'yours', 'yourself', '
yourselves', 'he', 'him', 'his', 'himself', \
        'she', "she's", 'her', 'hers', 'herself', 'it', "
it's", 'its', 'itself', 'they', 'them', 'their', \
        'theirs', 'themselves', 'what', 'which', 'who', '
whom', 'this', 'that', "that'll", 'these', 'those', \
        'am', 'is', 'are', 'was', 'were', 'be', 'been', '
being', 'have', 'has', 'had', 'having', 'do', 'does', \
        'did', 'doing', 'a', 'an', 'the', 'and', 'but', '
if', 'or', 'because', 'as', 'until', 'while', 'of', \
        'at', 'by', 'for', 'with', 'about', 'against', 'b
etween', 'into', 'through', 'during', 'before', 'after', \
        'above', 'below', 'to', 'from', 'up', 'down', 'in
', 'out', 'on', 'off', 'over', 'under', 'again', 'further', \
        'then', 'once', 'here', 'there', 'when', 'where',
'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', \
        'most', 'other', 'some', 'such', 'only', 'own', '
same', 'so', 'than', 'too', 'very', \
        's', 't', 'can', 'will', 'just', 'don', "don't",
'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', \
        've', 'y', 'ain', 'aren', "aren't", 'couldn', "co
uldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', \
        "hadn't", 'hasn', "hasn't", 'haven', "haven't", '
isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', \
        "mustn't", 'needn', "needn't", 'shan', "shan't",
'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't",
\
        'won', "won't", 'wouldn', "wouldn't"]])

```

In [23]:

```

# Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)

```

```
sentence = BeautifulSoup(sentence, 'lxml').get_text()
sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())
```

```
100%|██████████| 87773/87773 [00:52<00:00, 165
7.56it/s]
```

In [0]:

```
preprocessed_reviews[1500]
```

Out[0]:

```
'wow far two two star reviews one obviously no
idea ordering wants crispy cookies hey sorry
reviews nobody good beyond reminding us look o
rdering chocolate oatmeal cookies not like com
bination not order type cookie find combo quit
e nice really oatmeal sort calms rich chocola
e flavor gives cookie sort coconut type consis
tency let also remember tastes differ given op
inion soft chewy cookies advertised not crispy
cookies blurb would say crispy rather chewy h
appen like raw cookie dough however not see ta
ste like raw cookie dough soft however confusi
on yes stick together soft cookies tend not in
dividually wrapped would add cost oh yeah choc
olate chip cookies tend somewhat sweet want so
mething hard crisp suggest nabiso ginger snaps
want cookie soft chewy tastes like combinatio
n chocolate oatmeal give try place second orde
r'
```

In [24]:

```
#Random splitting of data to separate train, cv and test data
```

```
Y = final['Score'].values
```

```
X = preprocessed_reviews
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33) # this is random splitting
```

```
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33) # this is random splitting
```

```
print(len(X_train), y_train.shape)
```

```
print(len(X_cv), y_cv.shape)
```

```
print(len(X_test), y_test.shape)
```

```
39400 (39400,)
```

```
19407 (19407,)
```

```
28966 (28966,)
```



## [3.2] Preprocessing Review Summary

In [0]:

```
## Similarly you can do preprocessing for review summary als  
0.
```

## **[4] Featurization**

## [4.1] BAG OF WORDS

In [25]:

```
#Bow
count_vect = CountVectorizer(min_df=10)
#count_vect = CountVectorizer(min_df=10, max_features=500) #
for kd tree
#in scikit-learn
count_vect.fit(X_train)
print("some feature names ", count_vect.get_feature_names()[:
10])
print('='*50)

final_counts_train = count_vect.transform(X_train)
final_counts_cv = count_vect.transform(X_cv)
final_counts_test = count_vect.transform(X_test)
#print("the type of count vectorizer ",type(final_counts))
#print("the shape of out text BOW vectorizer ",final_counts.g
et_shape())
print("the number of unique words ", final_counts_train.get_s
hape()[1])
```

```
some feature names  ['ability', 'able', 'absen
ce', 'absent', 'absolute', 'absolutely', 'abso
lutly', 'absorb', 'absorbed', 'absorption']
```

```
=====
====
```

```
the number of unique words  7739
```

## [4.2] Bi-Grams and n-Grams.

In [0]:

```
#bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html

# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

## [4.3] TF-IDF

In [26]:

```
#tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10,
max_features=500) # for kd tree
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(X_train)
print("some sample features(unique words in the corpus)",tf_i
df_vect.get_feature_names()[0:10])
print('='*50)
```

```
final_tf_idf_train = tf_idf_vect.transform(X_train)
final_tf_idf_cv = tf_idf_vect.transform(X_cv)
final_tf_idf_test = tf_idf_vect.transform(X_test)

#print("the type of count vectorizer ",type(final_tf_idf))
#print("the shape of out text TFIDF vectorizer ",final_tf_idf
.get_shape())
#print("the number of unique words including both unigrams an
d bigrams ", final_tf_idf.get_shape()[1])
```

```
some sample features(unique words in the corpu
s) ['ability', 'able', 'able buy', 'able drink
', 'able eat', 'able enjoy', 'able figure', 'a
ble find', 'able finish', 'able get']
```

```
=====
====
```

## [4.4] Word2Vec

In [27]:

```
# Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())
```

In [28]:

```
# Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >
12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word]
as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
```

```

want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred at least
    # 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50,
workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNe
ws-vectors-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep wa
nt_to_train_w2v = True, to train your own w2v ")

```

```

[('fantastic', 0.8449681997299194), ('awesome'
, 0.83317631483078), ('good', 0.82166898250579
83), ('excellent', 0.7990680932998657), ('terr
ific', 0.7915112972259521), ('perfect', 0.7635
703682899475), ('wonderful', 0.763299226760864
3), ('amazing', 0.7176131010055542), ('nice',
0.710258960723877), ('decent', 0.7000181674957
275)]

```

```

=====
====

```

```

[('greatest', 0.8191127777099609), ('best', 0.
7109125852584839), ('tastiest', 0.699251055717
4683), ('disgusting', 0.6615686416625977), ('n
astiest', 0.6208571791648865), ('smoothest', 0
.6163804531097412), ('terrible', 0.59675610065

```

```
4602), ('awful', 0.5936288833618164), ('hottest', 0.5926538705825806), ('wins', 0.5878374576568604)]
```

In [29]:

```
w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times

17386

sample words ['dogs', 'loves', 'chicken', 'product', 'china', 'wont', 'buying', 'anymore', 'hard', 'find', 'products', 'made', 'usa', 'one', 'isnt', 'bad', 'good', 'take', 'chances', 'till', 'know', 'going', 'imports', 'love', 'saw', 'pet', 'store', 'tag', 'attached', 'regarding', 'satisfied', 'safe', 'infestation', 'literally', 'everywhere', 'flying', 'around', 'kitchen', 'bought', 'hoping', 'least', 'get', 'rid', 'weeks', 'fly', 'stuck', 'squishing', 'buggers', 'success', 'rate']



## [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

### [4.4.1.1] Avg W2v

In [30]:

```
# average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is
                    # stored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50,
                             # you might need to change this to 300 if you
                             # use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_model:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))
```

```
100%|██████████| 87773/87773 [06:35<00:00, 221.80it/s]
```

87773

50

#### [4.4.1.2] TFIDF weighted W2v

In [31]:

```
# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the
idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

In [0]:

```
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
```

```
rd)]  
    # to reduce the computation we are  
    # dictionary[word] = idf value of word in whole c  
ourpus  
    # sent.count(word) = tf valeus of word in this re  
view  
    tf_idf = dictionary[word]*(sent.count(word)/len(s  
ent))  
    sent_vec += (vec * tf_idf)  
    weight_sum += tf_idf  
    if weight_sum != 0:  
        sent_vec /= weight_sum  
    tfidf_sent_vectors.append(sent_vec)  
    row += 1
```

```
100%|███████████████████████████████████████████████████████████  
██████████████████████████████████████████████████████████████| 4986/498  
6 [00:20<00:00, 245.63it/s]
```

# [5] Assignment 4: Apply Naive Bayes

## 1. Apply Multinomial NaiveBayes on these feature sets

- **SET 1:** Review text, preprocessed one converted into vectors using (BOW)
- **SET 2:** Review text, preprocessed one converted into vectors using (TFIDF)

## 2. The hyper paramter tuning(find best Alpha)

- Find the best hyper parameter which will give the maximum [AUC](#) value
- Consider a wide range of alpha values for hyperparameter tuning, start as low as 0.00001
- Find the best hyper paramter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

## 3. Feature importance


- Find the top 10 features of positive class and top 10 features of negative class for both feature sets **Set 1** and **Set 2** using values of `feature\_log\_prob\_` parameter of [MultinomialNB](#) and print their corresponding feature names

## 4. Feature engineering


- To increase the performance of your model, you can also experiment with with feature engineering like :
  - Taking length of reviews as another feature.
  - Considering some features from review summary as well.

## 5. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure. Here on X-axis you will have alpha values, since they have a wide range, just to represent those alpha values on the graph, apply log function on those alpha values.




- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.




- Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points. Please visualize your confusion matrices

using [seaborn heatmaps](#).



## 6. [Conclusion](#)

- [You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library link](#)



### **Note: Data Leakage**

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on you train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this [link](#).

# Applying Multinomial Naive Bayes



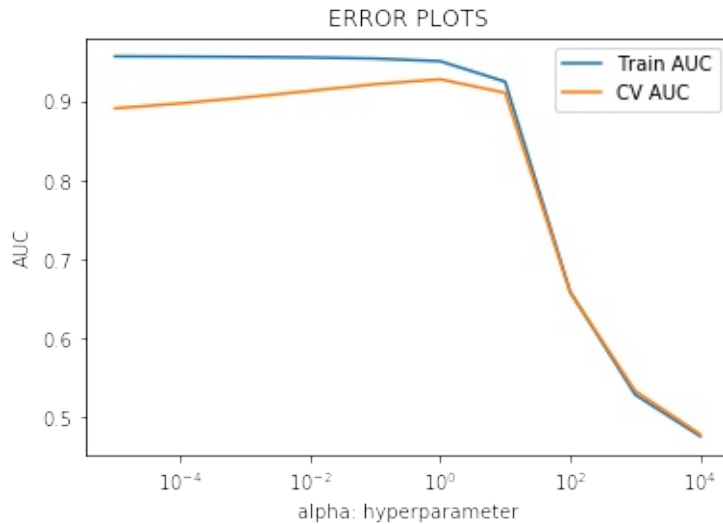
## [5.1] Applying Naive Bayes on BOW, SET 1

In [32]:

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import roc_auc_score
import matplotlib.pyplot as plt

train_auc = []
cv_auc = []
alpha_p = [0.00001, 0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0, 10000.0]
for i in alpha_p:
    clf = MultinomialNB(alpha=i)
    clf.fit(final_counts_train, y_train)
    y_train_pred = clf.predict_proba(final_counts_train)[:,-1]
    y_cv_pred = clf.predict_proba(final_counts_cv)[:,-1]
    train_auc.append(roc_auc_score(y_train, y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))

plt.semilogx(alpha_p, train_auc, label='Train AUC')
plt.semilogx(alpha_p, cv_auc, label='CV AUC')
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```

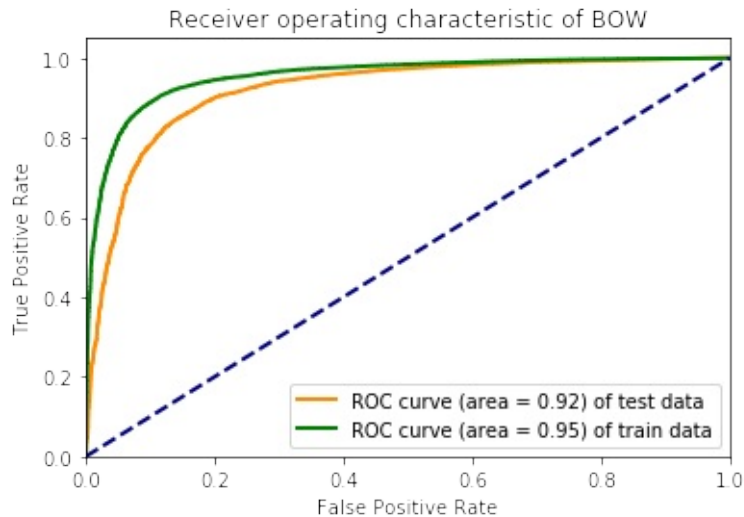


In [34]:

```
from sklearn.metrics import roc_curve, auc
clf = MultinomialNB(alpha=0.1)
clf.fit(final_counts_train, y_train)

y_train_pred = clf.predict_proba(final_counts_train)[: ,1]
y_test_pred = clf.predict_proba(final_counts_test)[: ,1]
fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
roc_auc_train = auc(fpr_train, tpr_train)
fpr_test, tpr_test, _ = roc_curve(y_test, y_test_pred)
roc_auc_test = auc(fpr_test, tpr_test)
plt.figure()
lw = 2
plt.plot(fpr_test, tpr_test, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f) of test data'
         % roc_auc_test)
plt.plot(fpr_train, tpr_train, color='green',
         lw=lw, label='ROC curve (area = %0.2f) of train data'
         % roc_auc_train)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of BOW')
plt.legend(loc="lower right")
plt.show()
```



In [35]:

```
from sklearn.metrics import confusion_matrix
from sklearn.utils.multiclass import unique_labels

clf = MultinomialNB(alpha=0.1)
clf.fit(final_counts_train, y_train)

y_test_pred = clf.predict(final_counts_test)
class_names = 0,1

def plot_confusion_matrix(y_true, y_pred, classes,
                           normalize=False,
                           title=None,
                           cmap=plt.cm.Blues):
    cm = confusion_matrix(y_true, y_pred)
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    s]
```

```

# Only use the labels that appear in the data
#classes = classes[unique_labels(y_true, y_pred)]
fig, ax = plt.subplots()
im = ax.imshow(cm, interpolation='nearest', cmap=cmap)
ax.figure.colorbar(im, ax=ax)
# We want to show all ticks...
ax.set(xticks=np.arange(cm.shape[1]),
       yticks=np.arange(cm.shape[0]),
       # ... and label them with the respective list entries
       xticklabels=classes, yticklabels=classes,
       title=title,
       ylabel='True label',
       xlabel='Predicted label')

# Rotate the tick labels and set their alignment.
plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
         rotation_mode="anchor")

# Loop over data dimensions and create text annotations.
fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(j, i, format(cm[i, j], fmt),
                ha="center", va="center",
                color="white" if cm[i, j] > thresh else "
black")
fig.tight_layout()
return ax

np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plot_confusion_matrix(y_test, y_test_pred, classes=class_names,
                      normalize=False,

```

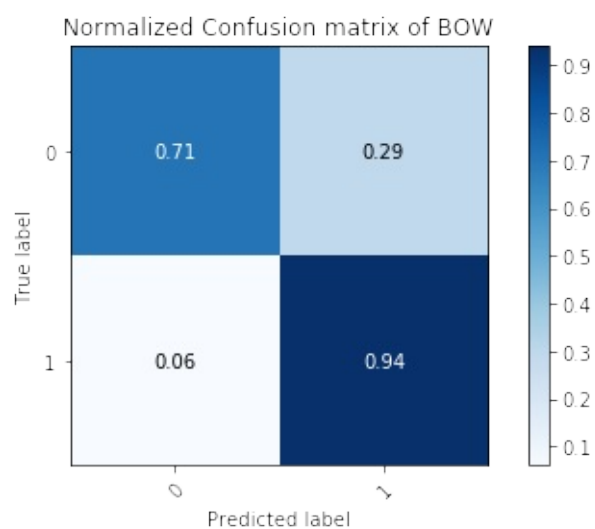

```
title='Confusion matrix of BOW')
```

```
plot_confusion_matrix(y_test, y_test_pred, classes=class_names,  
normalize=True,
```

```
title='Normalized Confusion matrix of BOW')
```

Out[35]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x2067c814908>



### [5.1.1] Top 10 important features of positive class from SET 1

In [36]:

```
# Please write all the code with proper documentation
import numpy as np
clf = MultinomialNB(alpha=0.1)
clf.fit(final_counts_train, y_train)

positive_class_prob_sorted = clf.feature_log_prob_[1, :].argsort()[::-1]
print("The log probabilities of top 10 features of positive class", sorted(clf.feature_log_prob_[1, :], reverse=True)[0:10])
print(" ")
print("The top 10 features of positive class", np.take(count_vector.get_feature_names(), positive_class_prob_sorted[0:10]))
```

The log probabilities of top 10 features of positive class [-3.6657699192897795, -4.484634831770842, -4.621432053884652, -4.684158912564502, -4.8299662808686765, -4.919466944281439, -4.974899637542597, -4.998627419499165, -5.007035922009258, -5.007905530491808]

The top 10 features of positive class ['not' 'like' 'good' 'great' 'one' 'taste' 'coffee' 'love' 'flavor' 'would']

### [5.1.2] Top 10 important features of negative class from SET 1

In [37]:

```

# Please write all the code with proper documentation
import numpy as np
clf = MultinomialNB(alpha=0.1)
clf.fit(final_counts_train, y_train)

negative_class_prob_sorted = clf.feature_log_prob_[0, :].argsort()[::-1]

print("The log probabilities of top 10 features of negative class", sorted(clf.feature_log_prob_[0, :], reverse=True)[0:10])
print(" ")
print("The top 10 features of negative class", np.take(count_vect.get_feature_names(), negative_class_prob_sorted[0:10]))

```

The log probabilities of top 10 features of negative class [-3.231232178287584, -4.32954268549976, -4.58757981454618, -4.604566301776175, -4.665211872752491, -4.870015278859359, -4.99394239970529, -5.068088430287161, -5.092054588438736, -5.1481755412909544]

The top 10 features of negative class ['not' 'like' 'taste' 'would' 'product' 'one' 'coffee' 'good' 'no' 'flavor']

# Performing feature engineering (adding length of review) on BOW

In [101]:

```
from sklearn.base import BaseEstimator, TransformerMixin

class ItemSelector(BaseEstimator, TransformerMixin):
    """For data grouped by feature, select subset of data at
    a provided key.

    The data is expected to be stored in a 2D data structure,
    where the first
    index is over features and the second is over samples. i
    .e.

    >> len(data[key]) == n_samples

    Please note that this is the opposite convention to sciki
    t-learn feature
    matrixes (where the first index corresponds to sample).

    ItemSelector only requires that the collection implement
    getitem
    (data[key]). Examples include: a dict of lists, 2D numpy
    array, Pandas
    DataFrame, numpy record array, etc.

    >> data = {'a': [1, 5, 2, 5, 2, 8],
                'b': [9, 4, 1, 4, 1, 3]}
    >> ds = ItemSelector(key='a')
    >> data['a'] == ds.transform(data)
```



*ItemSelector is not designed to handle data grouped by sample. (e.g. a list of dicts). If your data is structured this way, consider a transformer along the lines of `sklearn.feature\_extraction.DictVectorizer`.*

*Parameters*

*-----*

*key : hashable, required*

*The key corresponding to the desired value in a mappable.*

*"""*

```
def __init__(self, key):
    self.key = key
```

```
def fit(self, x, y=None):
    return self
```

```
def transform(self, data_dict):
    return data_dict[self.key]
```

```
class SubjectBodyExtractor(BaseEstimator, TransformerMixin):
    from tqdm import tqdm
    def fit(self, x, y=None):
        return self

    def transform(self, data):
        features = np.recarray(shape=(len(data),),
                                dtype=[('Summary1', object), ('Text1', object)])
        i=0
        for sentence1 in tqdm(data['Text'].values):
            sentence1 = re.sub(r"http\S+", "", sentence1)
            sentence1 = BeautifulSoup(sentence1, 'lxml').get_text()
            sentence1 = decontracted(sentence1)
```

```

        sentence1 = re.sub("\S*\d\S*", "", sentence1).strip()

        sentence1 = re.sub('[^A-Za-z]+', ' ', sentence1)
        # https://gist.github.com/sebleier/554280
        sentence1 = ' '.join(e.lower() for e in sentence1
        .split() if e.lower() not in stopwords)
        features['Text1'][i] = sentence1.strip()
        i=i+1
    i=0
    for sentence11 in tqdm(data['Summary'].values):
        sentence11 = re.sub(r"http\S+", "", sentence11)
        sentence11 = BeautifulSoup(sentence11, 'lxml').get_text()

        sentence11 = decontracted(sentence11)
        sentence11 = re.sub("\S*\d\S*", "", sentence11).strip()

        sentence11 = re.sub('[^A-Za-z]+', ' ', sentence11
        )
        # https://gist.github.com/sebleier/554280
        sentence11 = ' '.join(e.lower() for e in sentence
        11.split() if e.lower() not in stopwords)
        features['Summary1'][i] = sentence11.strip()
        i=i+1

    return features

```

In [102]:

```

#print(X_train1.shape)
X_train1 = timesorted_data[:28965]
y_train1 = X_train1['Score'].values
X_cv1 = timesorted_data[28965:57930]
X_test1 = timesorted_data[57930:]
y_test1 = X_test1['Score'].values
print(y_train1.shape)

```

```
#print(y_train1)
```

```
(28965,)
```

In [103]:

```
from sklearn.metrics import roc_curve, auc
from sklearn.pipeline import Pipeline, FeatureUnion
from sklearn.preprocessing import FunctionTransformer
from sklearn.metrics import roc_auc_score
import matplotlib.pyplot as plt

#adding a function to find the length of the review text
def get_text_length(x):
    return np.array([len(t) for t in x]).reshape(-1, 1)

# Featureunion adding the length of the text to the BOW vectors. Hyperparameter considered alpha=0.1
union = Pipeline([
    ('subjectbody', SubjectBodyExtractor()),
    ('features', FeatureUnion([
        ('summarytext', Pipeline([
            ('selector', ItemSelector(key='Summary1')),
            ('vectorizer', CountVectorizer(min_df=10)),
        ])),
        ('text', Pipeline([
            ('selector', ItemSelector(key='Text1')),
            ('vectorizer', CountVectorizer(min_df=10)),
        ])),
        ('length', Pipeline([
            ('selector', ItemSelector(key='Text1')),
            ('count', FunctionTransformer(get_text_length, validate=False)),
        ]))
    ])),
    ('clf', (MultinomialNB(alpha=0.1))))]
```

```

#fitting only to train data to avoid data leakage
features_train = union.fit(X_train1,y_train1)
y_train_pred = union.predict_proba(X_train1)[: ,1]
y_test_pred = union.predict_proba(X_test1)[: ,1]

fpr_train, tpr_train, _ = roc_curve(y_train1, y_train_pred)
roc_auc_train = auc(fpr_train, tpr_train)
fpr_test, tpr_test, _ = roc_curve(y_test1, y_test_pred)
roc_auc_test = auc(fpr_test, tpr_test)
plt.figure()
lw = 2
plt.plot(fpr_test, tpr_test, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f) of test data'
         % roc_auc_test)
plt.plot(fpr_train, tpr_train, color='green',
         lw=lw, label='ROC curve (area = %0.2f) of train data'
         % roc_auc_train)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of BOW')
plt.legend(loc="lower right")
plt.show()

y_test_prediction = union.predict(X_test1)
class_names = 0,1

# Plot non-normalized confusion matrix
plot_confusion_matrix(y_test1, y_test_prediction, classes=cla
ss_names, normalize=False,
                     title='Confusion matrix of BOW')
plot_confusion_matrix(y_test1, y_test_prediction, classes=cla
ss_names, normalize=True,
                     title='Normalized Confusion matrix of B

```

OW')

100%|██████████| 28965/28965 [00:21<00:00, 1335.04it/s]

6%|█| 1749/28965 [00:00<00:13, 2077.36it/s]C:\Users\nairh\Anaconda3\lib\site-packages\bs4\\_\_init\_\_.py:272: UserWarning: "b'...' " looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.

' BeautifulSoup.' % markup)

100%|██████████| 28965/28965 [00:14<00:00, 1987.13it/s]

100%|██████████| 28965/28965 [00:20<00:00, 1424.43it/s]

6%|█| 1759/28965 [00:00<00:13, 2044.63it/s]C:\Users\nairh\Anaconda3\lib\site-packages\bs4\\_\_init\_\_.py:272: UserWarning: "b'...' " looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.

' BeautifulSoup.' % markup)


100%|██████████| 28965/28965 [00:13<00:00, 2147.37it/s]

100%|██████████| 29843/29843 [00:21<00:00, 1358.01it/s]

49%|██████| 14612/29843 [00:07<00:06, 2184.60it/s]C:\Users\nairh\Anaconda3\lib\site-packages\bs4\\_\_init\_\_.py:272: UserWarning: "b'...' " looks like a filename, not markup. You should probably open this file and pass the filehandle into BeautifulSoup.

' BeautifulSoup.' % markup)


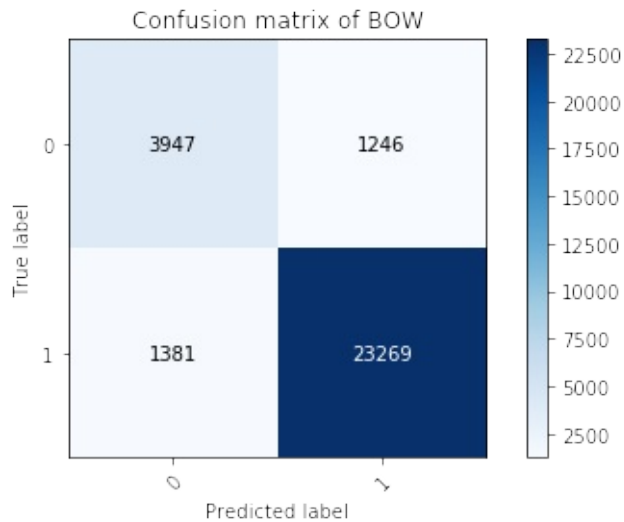
100%|██████████| 29843/29843 [00:15<00:00, 1948.33it/s]



```
100%|██████████| 29843/29843 [00:22<00:00, 134
3.41it/s]
 49%|██████    | 14670/29843 [00:07<00:08, 168
7.04it/s]C:\Users\nairh\Anaconda3\lib\site-pac
kages\bs4\__init__.py:272: UserWarning: "b'...
'" looks like a filename, not markup. You shou
ld probably open this file and pass the fileha
ndle into BeautifulSoup.
  ' BeautifulSoup.' % markup)
100%|██████████| 29843/29843 [00:16<00:00, 176
7.93it/s]
```

Out[103]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x20
602c18748>
```



Feature engineering has been performed by incorporating the length of the review text and the review summary in addition to the usual review texts. Here, data has been split after sorting the data based on time. Inclusion of these features has improved the AUC from 0.92 to 0.94. This improvement in result can also be observed in the increase of the diagonal terms in confusion matrix as the TN increases from 0.71 to 0.76.

## [5.2] Applying Naive Bayes on TFIDF, SET 2


In [38]:

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import roc_auc_score
import matplotlib.pyplot as plt

train_auc = []
cv_auc = []
alpha_p = [0.00001, 0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0, 10000.0]
for i in alpha_p:
    clf = MultinomialNB(alpha=i)
    clf.fit(final_tf_idf_train, y_train)
    y_train_pred = clf.predict_proba(final_tf_idf_train)[:,-1]
    y_cv_pred = clf.predict_proba(final_tf_idf_cv)[:,-1]
    train_auc.append(roc_auc_score(y_train, y_train_pred))
    cv_auc.append(roc_auc_score(y_cv, y_cv_pred))

plt.semilogx(alpha_p, train_auc, label='Train AUC')
plt.semilogx(alpha_p, cv_auc, label='CV AUC')
plt.legend()
plt.xlabel("alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.show()
```






In [39]:

```
from sklearn.metrics import roc_curve, auc
clf = MultinomialNB(alpha=0.1)
clf.fit(final_tf_idf_train, y_train)

y_train_pred = clf.predict_proba(final_tf_idf_train)[:,-1]
y_test_pred = clf.predict_proba(final_tf_idf_test)[:,-1]
fpr_train, tpr_train, _ = roc_curve(y_train, y_train_pred)
roc_auc_train = auc(fpr_train, tpr_train)
fpr_test, tpr_test, _ = roc_curve(y_test, y_test_pred)
roc_auc_test = auc(fpr_test, tpr_test)
plt.figure()
lw = 2
plt.plot(fpr_test, tpr_test, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f) of test data'
         % roc_auc_test)
plt.plot(fpr_train, tpr_train, color='green',
         lw=lw, label='ROC curve (area = %0.2f) of train data'
         % roc_auc_train)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of BOW')
plt.legend(loc="lower right")
plt.show()
```



In [40]:

```
clf = MultinomialNB(alpha=0.1)
clf.fit(final_tf_idf_train, y_train)

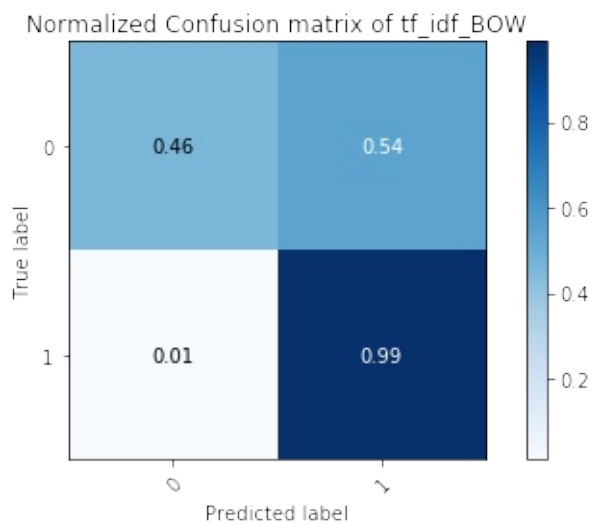
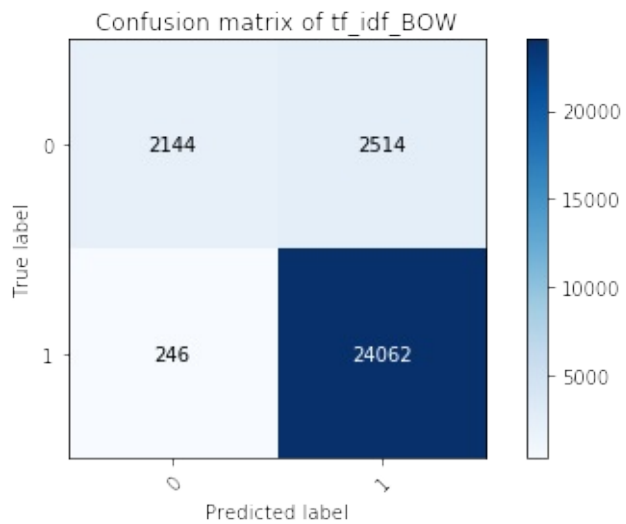
y_test_pred = clf.predict(final_tf_idf_test)
class_names = 0,1

# Plot non-normalized confusion matrix
plot_confusion_matrix(y_test, y_test_pred, classes=class_names,
                      normalize=False,
                      title='Confusion matrix of tf_idf_BOW')
plot_confusion_matrix(y_test, y_test_pred, classes=class_names,
                      normalize=True,
                      title='Normalized Confusion matrix of tf_idf_BOW')
```

Out[40]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x20

6005735f8>



### [5.2.1] Top 10 important features of positive class from SET 2

In [41]:

```
# Please write all the code with proper documentation
import numpy as np
clf = MultinomialNB(alpha=0.1)
clf.fit(final_tf_idf_train, y_train)
```

```
positive_class_prob_sorted = sorted(clf.feature_log_prob_[1,
:], reverse=True)

print("The log probabilities of top 10 features of positive c
lass",positive_class_prob_sorted[0:10])
```

The log probabilities of top 10 features of po  
sitive class [-5.216932987431211, -5.558904724  
233968, -5.635538469939925, -5.685468623843244  
, -5.742868367815882, -5.786356067421885, -5.8  
17105572467525, -5.91497373459664, -5.92830528  
42277, -5.930149922811605]

## [5.2.2] Top 10 important features of negative class from SET 2

In [42]:

```
# Please write all the code with proper documentation
import numpy as np
clf = MultinomialNB(alpha=0.1)
clf.fit(final_tf_idf_train, y_train)

negative_class_prob_sorted = sorted(clf.feature_log_prob_[0,
:], reverse=True)

print( "The log probabilities of top 10 features of negative
class",negative_class_prob_sorted[:10])
```

The log probabilities of top 10 features of ne  
gative class [-4.70841414194601, -5.4894960930  
3242, -5.609663505693307, -5.625675153849732,  
-5.654216648135872, -5.79220658441217, -5.9753  
67081149988, -6.048529356022203, -6.0886905628  
8089, -6.170275329601554]



## [6] Conclusions

In [47]:

```
# Please compare all your models using Prettytable library
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Model", "Hyperparameter", "AUC"]

x.add_row(["BOW", "MultinomialNB", 0.1, 0.92])
x.add_row(["tfidf BOW", "MultinomialNB", 0.1, 0.95])

print(x)
```

```
+-----+-----+-----+
+-----+
| Vectorizer |      Model      | Hyperparameter
| AUC  |
+-----+-----+-----+
+-----+
|      BOW      | MultinomialNB |      0.1
| 0.92 |
| tfidf BOW | MultinomialNB |      0.1
| 0.95 |
+-----+-----+-----+
+-----+
```